

# DATA MINING

## LECTURE 10

---

### **Classification**

Basic Concepts

Decision Trees

# Catching tax-evasion

<i>Tid</i>	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Tax-return data for year 2011

A new tax return for 2012

Is this a cheating tax return?

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

An instance of the classification problem: learn a method for discriminating between records of different **classes** (**cheaters** vs **non-cheaters**)

# What is classification?

$$x \rightarrow \boxed{\checkmark \neq} \rightarrow y, y = \{ \text{yes, No} \}$$

- Classification** is the task of **learning a target function  $f$**  that maps attribute set  $x$  to one of the predefined class labels  $y$

$$x = \{ \text{Refund, Marital Status, Taxable Income} \}$$

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

One of the attributes is the **class attribute**  
In this case: Cheat

Two **class labels** (or **classes**): Yes (1), No (0)




Figure 4.2. Classification as the task of mapping an input attribute set  $x$  into its class label  $y$ .

# Why classification?

- The target function  $f$  is known as a **classification model**
- **Descriptive modeling:** **Explanatory tool** to distinguish between objects of different classes (e.g., understand why people cheat on their taxes)
- **Predictive modeling:** Predict a class of a **previously unseen** record

# Examples of Classification Tasks

- Predicting tumor cells as benign or malignant
- Classifying credit card transactions as legitimate or fraudulent
- Categorizing news stories as finance, weather, entertainment, sports, etc
- Identifying spam email, spam web pages, adult content
- Understanding if a web query has commercial intent or not

# General approach to classification

training  
test

- Training set consists of records with known class labels
- Training set is used to build a classification model
- A labeled test set of previously unseen data records is used to evaluate the quality of the model.
- The classification model is applied to new records with unknown class labels

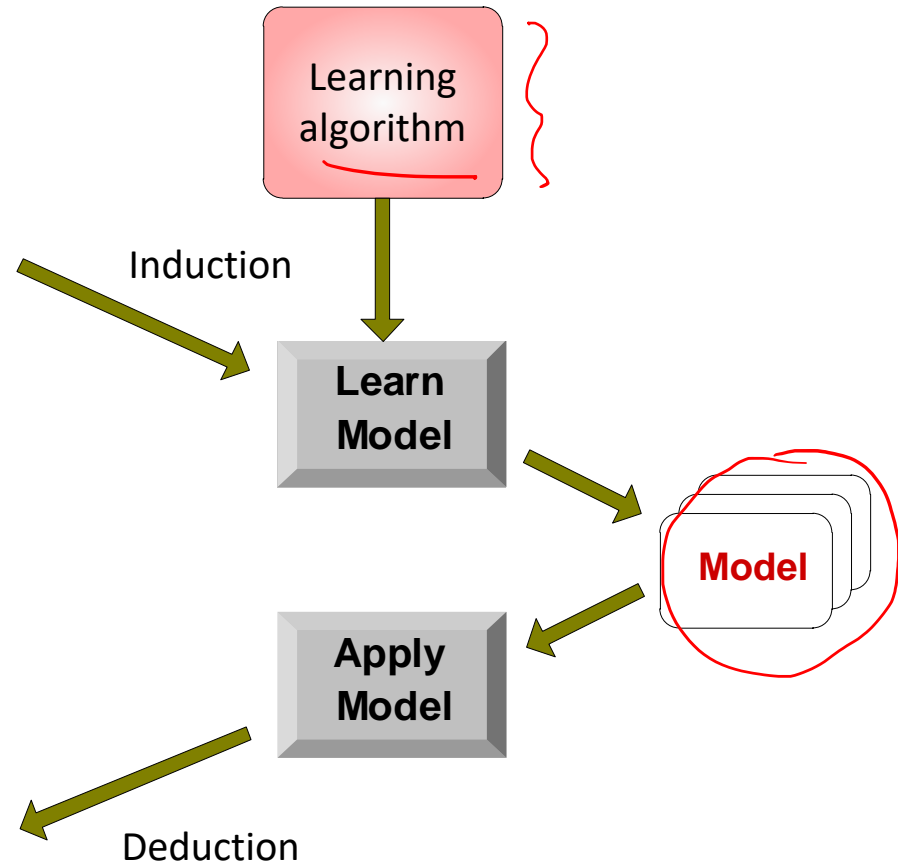
# Illustrating Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Total no. of ~~the~~ records:   
 FN  $\rightarrow$  are the no. of <sup>True</sup> records missed by the model to predict correctly.   
 TP  $\rightarrow$  are the no. of <sup>True</sup> records correctly predicted by the model.

# Evaluation of classification models

- Counts of **test records** that are correctly (or incorrectly) predicted by the classification model
- Confusion matrix** *Actual Class*

	Predicted Class	
	Class = 1	Class = 0
Actual Class		
Class = 1	$f_{11}$ <b>TP</b>	$f_{10}$ <b>FN</b>
Class = 0	$f_{01}$ <b>FP</b>	$f_{00}$ <b>TN</b>

Precision Recall F score   
 Accuracy =  $\frac{\text{\# correct predictions}}{\text{total \# of predictions}} = \frac{f_{11} + f_{00}}{f_{11} + f_{10} + f_{01} + f_{00}}$

Error rate =  $\frac{\text{\# wrong predictions}}{\text{total \# of predictions}} = \frac{f_{10} + f_{01}}{f_{11} + f_{10} + f_{01} + f_{00}}$

$P = \frac{TP + FP}{TP + FP + FN}$    
 $R = \frac{TP + FN}{TP + FN + FP}$    
 $F1 = \frac{2 * P * R}{P + R}$



# Classification Techniques

- Decision Tree based Methods
- Rule-based Methods
- Memory based reasoning
- Neural Networks
- Naïve Bayes and Bayesian Belief Networks
- Support Vector Machines


# Classification Techniques

- Decision Tree based Methods
- Rule-based Methods
- Memory based reasoning
- Neural Networks
- Naïve Bayes and Bayesian Belief Networks
- Support Vector Machines

# Decision Trees

## ! Decision tree

- A **flow-chart-like tree** structure
- **Internal node** denotes a **test on an attribute**
- **Branch** represents an **outcome of the test**
- **Leaf nodes** represent **class labels** or class distribution



Root node →

Internal nodes

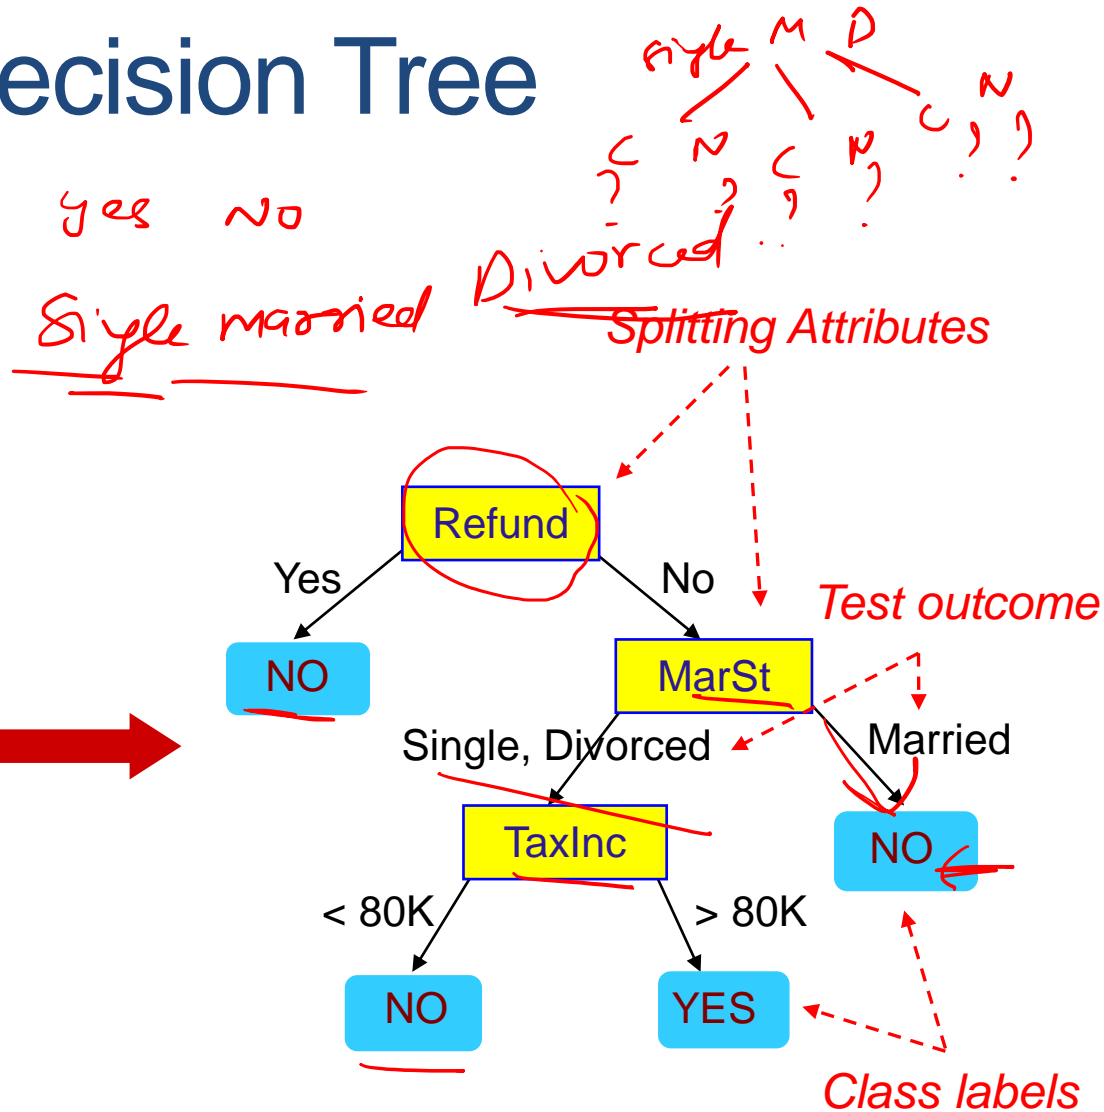
Leaf or terminal nodes

# Example of a Decision Tree

categorical  
categorical  
continuous  
class

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Training Data

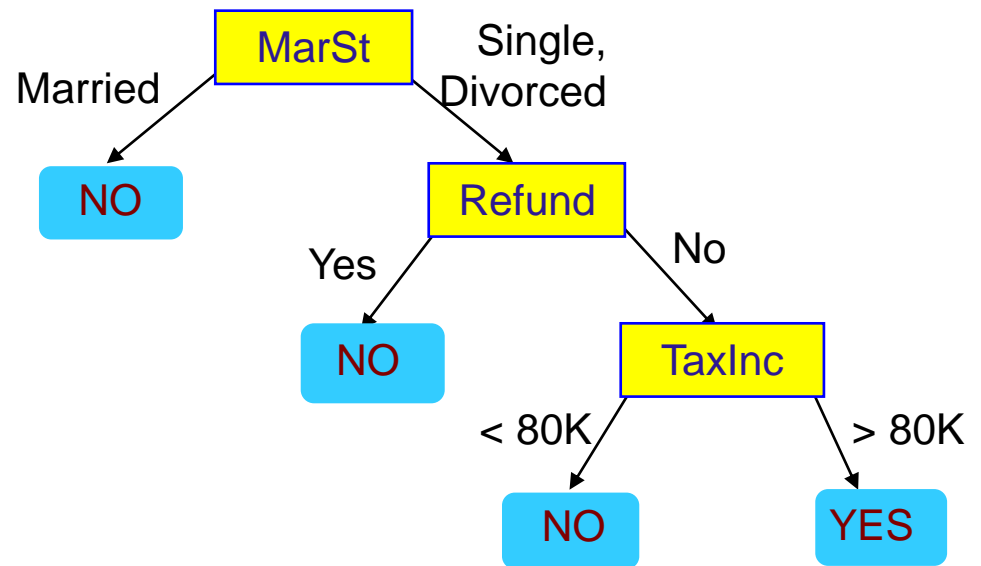


Model: Decision Tree

# Another Example of Decision Tree

*categorical*  
*categorical*  
*continuous*  
*class*

<i>Tid</i>	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



There could be more than one tree that fits the same data!


# Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

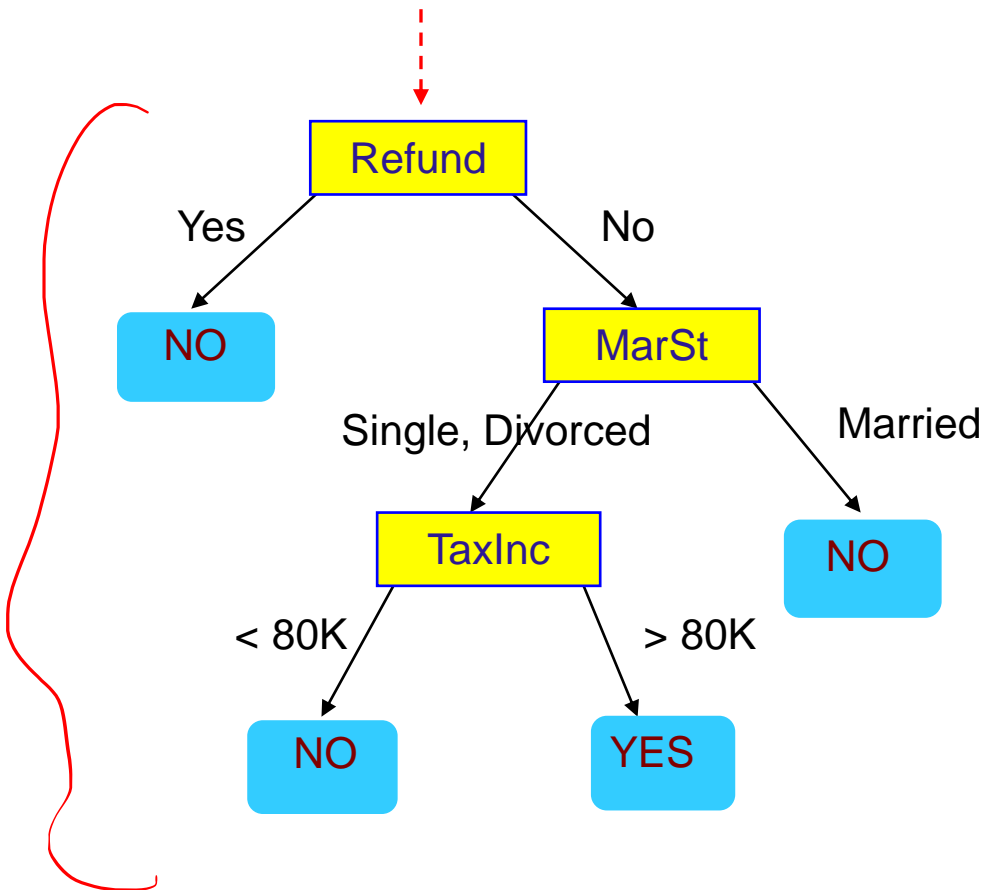
Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



# Apply Model to Test Data

Start from the root of tree.



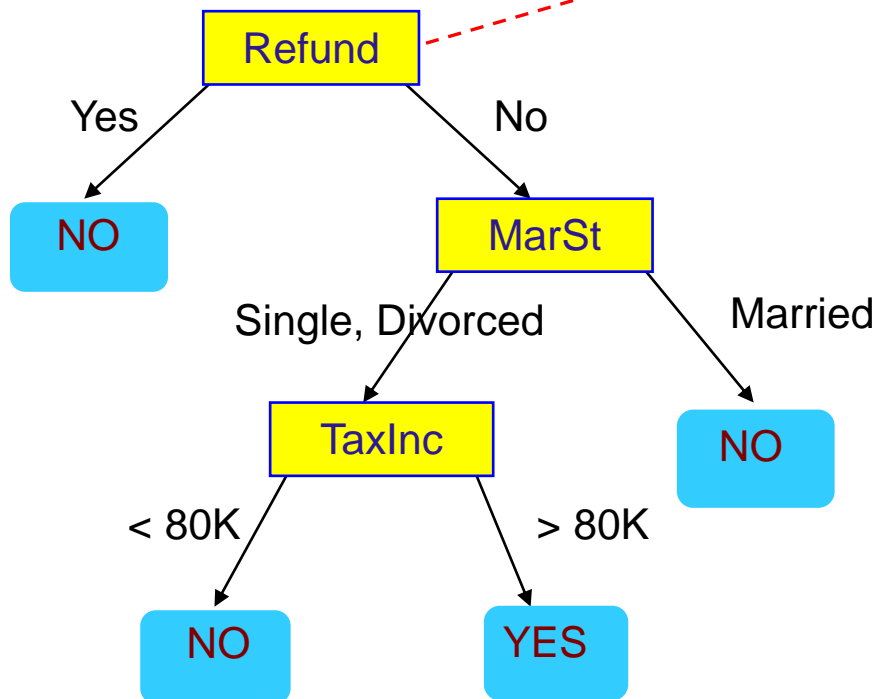
Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

# Apply Model to Test Data

Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

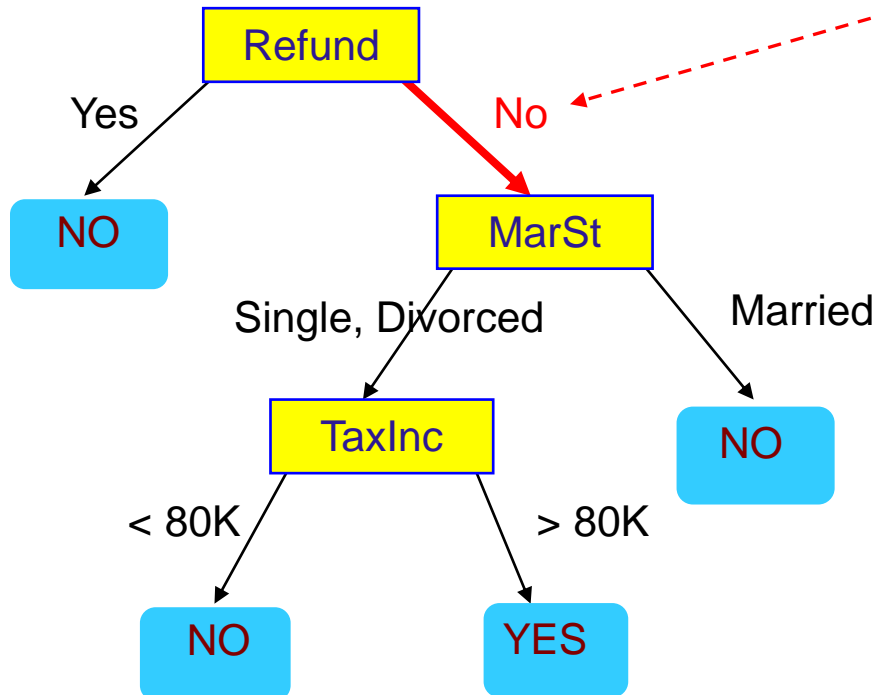




# Apply Model to Test Data

Test Data

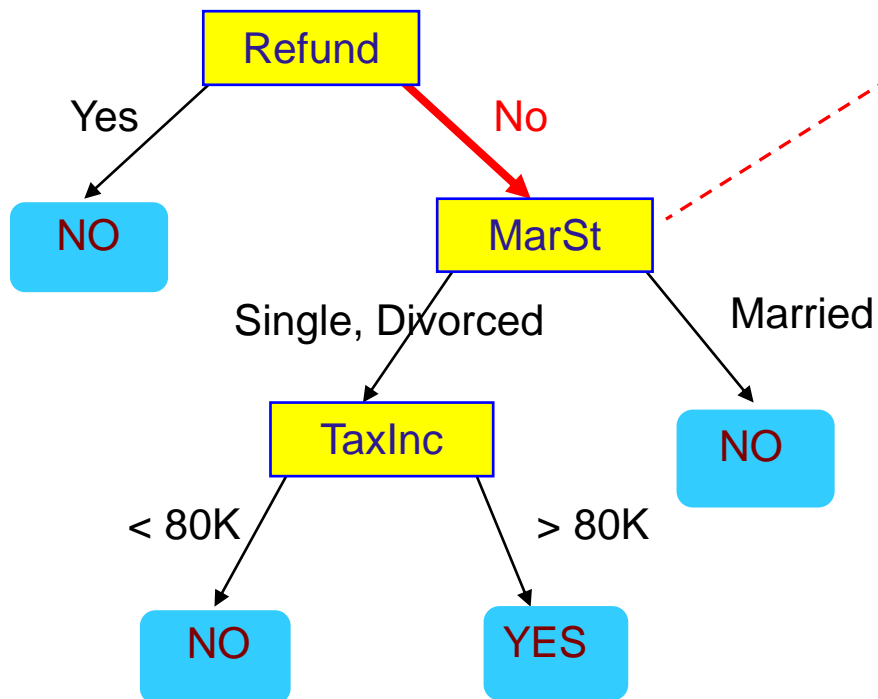
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



# Apply Model to Test Data

Test Data

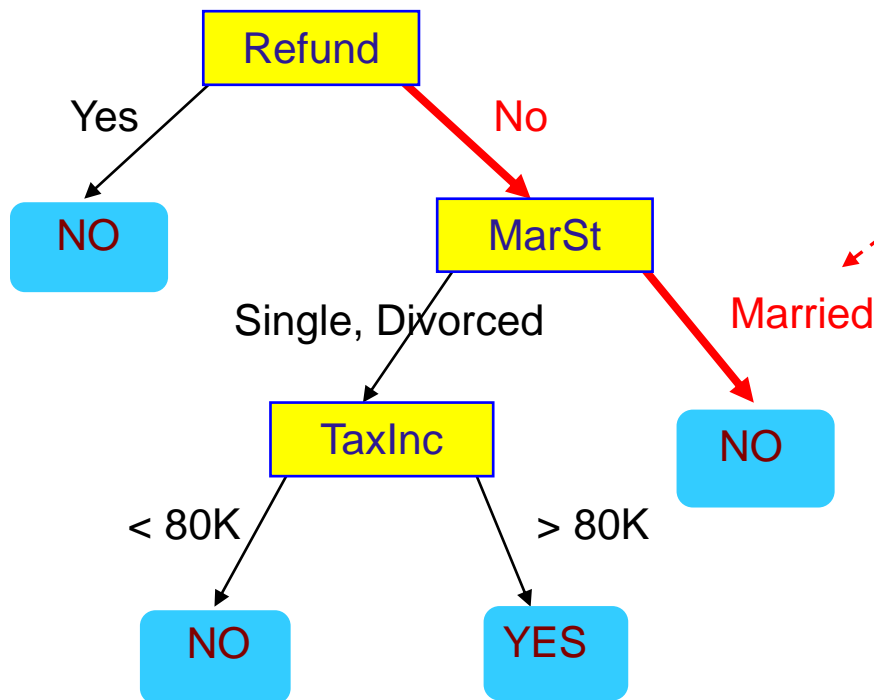
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



# Apply Model to Test Data

Test Data

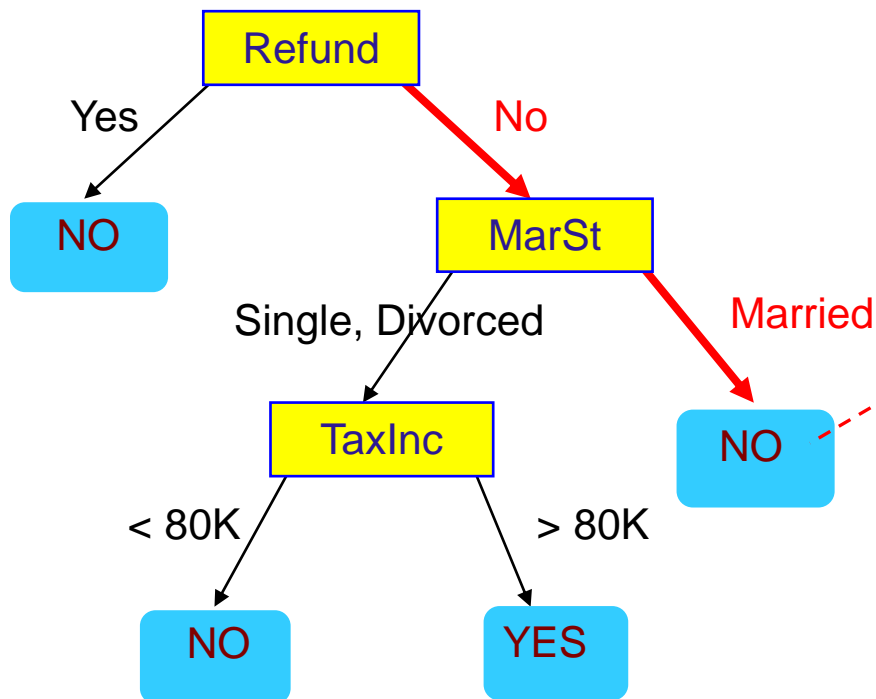
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



# Apply Model to Test Data

Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Assign Cheat to "No"

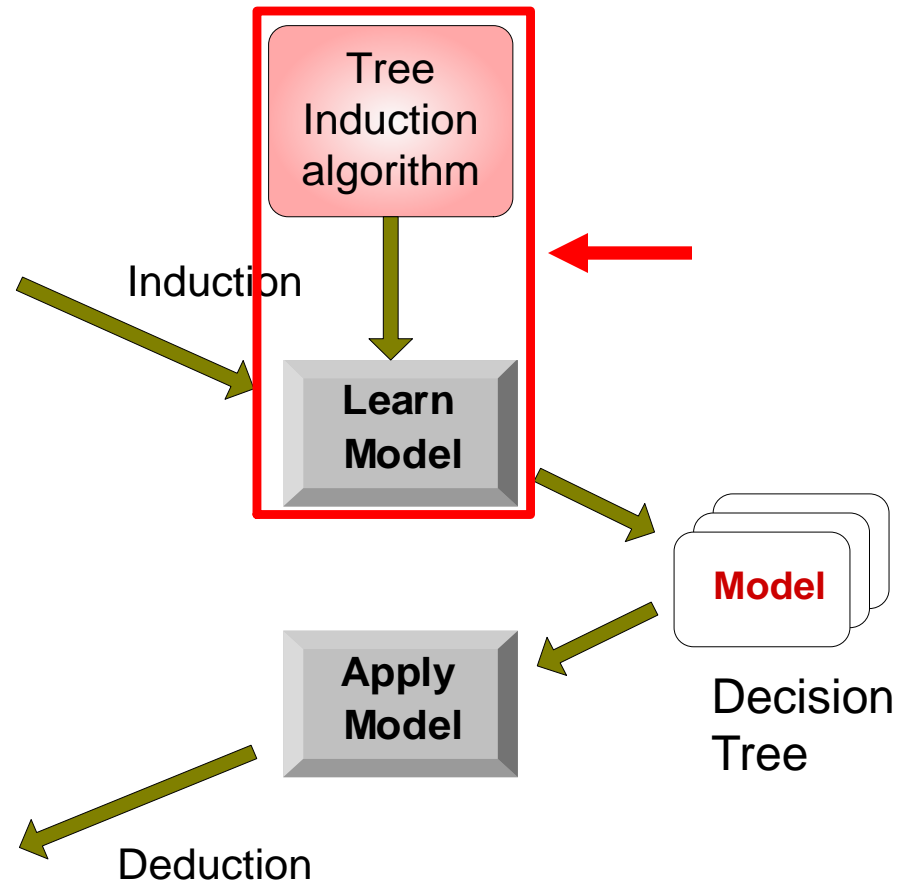
# Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



# Tree Induction

- Finding the best decision tree is **NP-hard**
- **Greedy** strategy.
  - Split the records based on an attribute test that optimizes **certain criterion**.
- Many Algorithms:
  - Hunt's Algorithm (one of the earliest)
  - CART
  - ID3, C4.5
  - SLIQ, SPRINT


Refund  $\rightarrow$  node  $t$

$D \rightarrow$  Dataset (table),  $D_t \rightarrow$  column 1  $D_{\text{Refund}} \rightarrow$  column 1

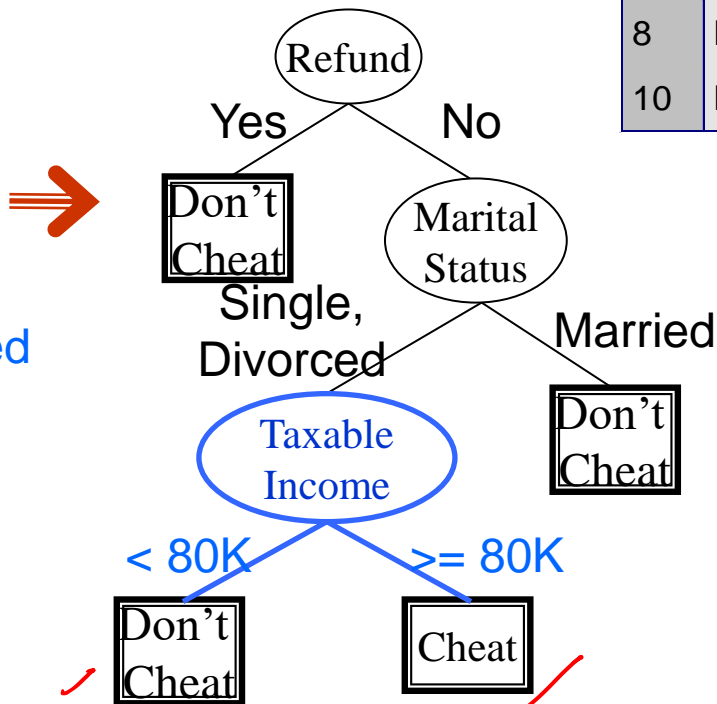
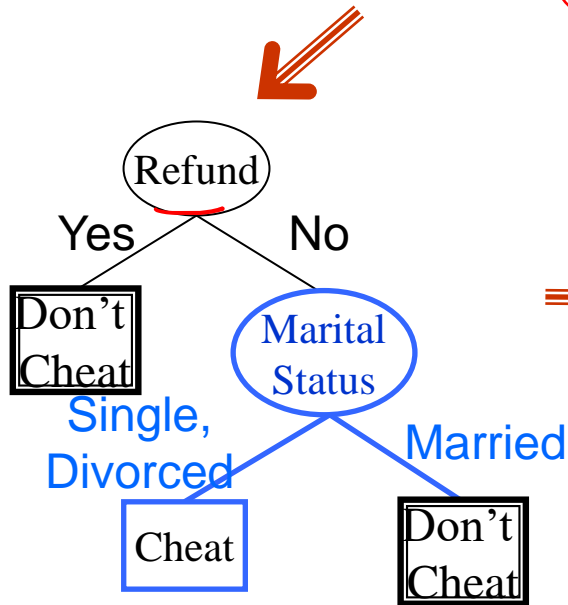

# General Structure of Hunt's Algorithm

- Let  $D_t$  be the set of training records that reach a node  $t$   
 $D_t \rightarrow D_{\text{Refund}}$
- General Procedure:
  - If  $D_t$  contains records that belong to the **same** class  $y_t$ , then  $t$  is a leaf node labeled as  $y_t$
  - If  $D_t$  contains records with the same attribute values, then  $t$  is a leaf node labeled with the **majority class**  $y_t$
  - If  $D_t$  is an **empty set**, then  $t$  is a leaf node labeled by the **default class**,  $y_d$
  - If  $D_t$  contains records that belong to more than one class, use an attribute test to split the data into smaller subsets.
    - Recursively apply the procedure to each subset.

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



# Hunt's Algorithm



<u>Tid</u>	<u>Refund</u>	<u>Marital Status</u>	<u>Taxable Income</u>	<u>Cheat</u>
1	Yes	Single	125K	No
4	Yes	Married	120K	No
7	Yes	Divorced	220K	No
2	No	Married	100K	No
6	No	Married	60K	No
9	No	Married	75K	No
3	No	Single	70K	No
5	No	Divorced	95K	Yes
8	No	Single	85K	Yes
10	No	Single	90K	Yes



# Constructing decision-trees (pseudocode)

**GenDecTree**(Sample **S**, Features **F**)

1. If **stopping\_condition**(**S**,**F**) = true then
  - a. leaf = **createNode**()
  - b. leaf.label = **Classify**(**S**)
  - c. return leaf
2. root = **createNode**()
3. root.test\_condition = **findBestSplit**(**S**,**F**)
4. **V** = {**v** | **v** a possible outcome of root.test\_condition}
5. for each value **v**  $\in$  **V**:
  - a. **S<sub>v</sub>** = {**s** | root.test\_condition(**s**) = **v** and **s**  $\in$  **S**};
  - b. child = **GenDecTree**(**S<sub>v</sub>**, **F**) ;
  - c. Add child as a descent of root and label the edge (root→child) as **v**
6. return root

# Tree Induction


- Issues
  - How to **Classify** a leaf node
    - Assign the **majority class**
    - If leaf is empty, assign the **default class** – the class that has the highest popularity.
  - Determine how to split the records
    - How to specify the attribute test condition?
    - How to determine the best split?
  - Determine when to stop splitting

# How to Specify Test Condition?


- Depends on attribute types
  - Nominal
  - Ordinal
  - Continuous
- Depends on number of ways to split
  - 2-way split
  - Multi-way split

# Splitting Based on Nominal Attributes


- **Multi-way split:** Use as many partitions as distinct values.



- **Binary split:** Divides values into two subsets.  
Need to find optimal partitioning.




OR




# Splitting Based on Ordinal Attributes


- **Multi-way split:** Use as many partitions as distinct values.




- **Binary split:** Divides values into two subsets – respects the order. Need to find optimal partitioning.



OR

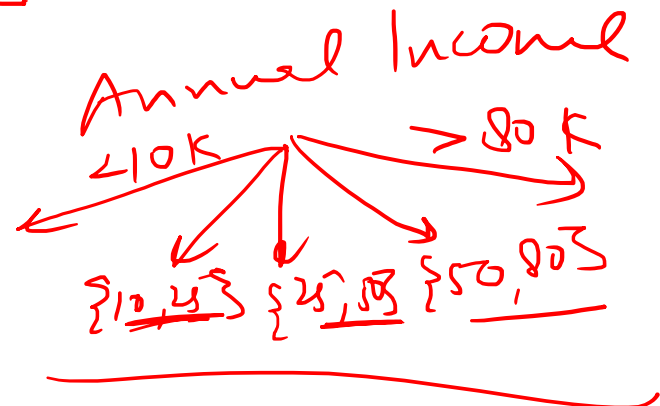



- What about this split?




# Splitting Based on Continuous Attributes


- Different ways of handling
  - **Discretization** to form an **ordinal** categorical attribute
    - **Static** – discretize once at the beginning
    - **Dynamic** – ranges can be found by equal interval bucketing, equal frequency bucketing (percentiles), or clustering.
  - **Binary Decision**:  $(A < v)$  or  $(A \geq v)$ 
    - consider all possible splits and finds the best cut
    - can be more compute intensive



# Splitting Based on Continuous Attributes



(i) Binary split

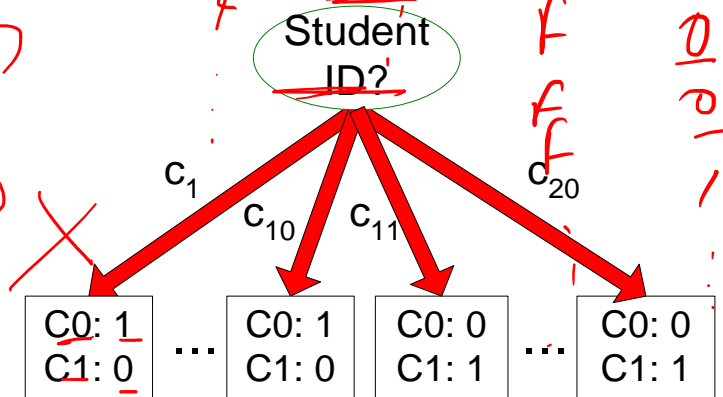




(ii) Multi-way split

# How to determine the Best Split

Before Splitting: 10 records of class 0,  
10 records of class 1

Std Id	<u>Own Car</u>	Car type	label
1	<u>Yes</u>	Family	0
2	<u>Yes</u>	S	0
3	<u>No</u>	L	1
4	<u>No</u>	F	0
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...



Which test condition is the best?


C0: 10  
C1: 0

C0: 0  
C1: 10




# How to determine the Best Split

- **Greedy** approach:
  - Nodes with **homogeneous** class distribution are preferred
- Need a measure of node **impurity**:



Non-homogeneous,  
High degree of impurity



Homogeneous,  
Low degree of impurity

C0: 10  
C1: 0

- Ideas?

# Measuring Node Impurity

- $p(i|t)$ : fraction of records associated with node  $t$  belonging to class  $i$

yes/no

Refund

$$\text{Entropy}(t) = - \sum_{i=1}^c p(i|t) \log p(i|t)$$

✓  $0 \log 0 = 0$   $p(0|t) = 0$

- Used in ID3 and C4.5

$$\text{Gini}(t) = 1 - \sum_{i=1}^c [p(i|t)]^2$$

✓

- Used in CART, SLIQ, SPRINT.

$$\text{Classification error}(t) = 1 - \max_i [p(i|t)]$$

✓

# Gain

$I(\cdot)$  is the impurity measure of the given node  
 $N$  is the total no. of records at the parent node  
 $k$  is the no. of att. values

- **Gain of an attribute split:** compare the impurity of the parent node with the average impurity of the child nodes

$$\Delta = I(\text{parent}) - \sum_{j=1}^k \frac{N(v_j)}{N} I(v_j)$$

no. of records associated with the child node  $v_j$

- Maximizing the gain  $\Leftrightarrow$  Minimizing the weighted average impurity measure of children nodes
- If  $I() = \text{Entropy}()$ , then  $\Delta_{\text{info}}$  is called information gain

# Example

C1	<b>0</b>
C2	<b>6</b>

$$\begin{aligned}
 & \left. \begin{array}{l} C1 \quad 3 \\ C2 \quad 3 \end{array} \right\} \begin{aligned} & Gini = 1 - (3/6)^2 - (3/6)^2 = 0.5 \\ & Entropy = -3/6 \log(3/6) - (3/6) \log(3/6) \\ & Error = 1 - \max[3/6, 3/6] = 0.5 = 1 \end{aligned}
 \end{aligned}$$

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Gini = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = 0$$

$$Entropy = -0 \log 0 - 1 \log 1 = -0 - 0 = 0$$

$$Error = 1 - \max(0, 1) = 1 - 1 = 0$$

$$1 - \sum_{i=1}^C P(C_i)^2$$

C1	<b>1</b>
C2	<b>5</b>

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Gini = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

$$Entropy = - (1/6) \log_2 (1/6) - (5/6) \log_2 (1/6) = 0.65$$

$$Error = 1 - \max(1/6, 5/6) = 1 - 5/6 = 1/6$$

C1	<b>2</b>
C2	<b>4</b>

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$Gini = 1 - (2/6)^2 - (4/6)^2 = 0.444$$

$$Entropy = - (2/6) \log_2 (2/6) - (4/6) \log_2 (4/6) = 0.92$$


$$Error = 1 - \max(2/6, 4/6) = 1 - 4/6 = 1/3$$

# Impurity measures

- All of the impurity measures take value zero (**minimum**) for the case of a pure node where a single value has probability 1
- All of the impurity measures take **maximum** value when the class distribution in a node is **uniform**.

# Comparison among Splitting Criteria

For a 2-class problem:



# Categorical Attributes

- For **binary** values split in two
- For **multivalued** attributes, for each distinct value, gather counts for each class in the dataset
  - Use the **count matrix** to make decisions

Multi-way split

	CarType		
	Family	Sports	Luxury
C1	1	2	1
C2	4	1	1
Gini	0.393		

Two-way split  
(find best partition of values)

	CarType	
	{Sports, Luxury}	{Family}
C1	3	1
C2	2	4
Gini	0.400	


	CarType	
	{Sports}	{Family, Luxury}
C1	2	2
C2	1	5
Gini	0.419	

# Continuous Attributes

- Use Binary Decisions based on one value
- Choices for the **splitting value**
  - Number of possible splitting values = Number of **distinct values**
- Each **splitting value** has a **count matrix** associated with it
  - Class counts in each of the partitions,  $A < v$  and  $A \geq v$
- **Exhaustive** method to choose best  $v$ 
  - For each  $v$ , scan the database to gather count matrix and compute the impurity index
  - Computationally Inefficient! Repetition of work.



Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes





# Continuous Attributes

$$Gini(\leq 55) = 1 - (0 + 0) = 1$$

$$Gini(> 55) = 1 - (3/10)^2 - (7/10)^2 = 0.42$$

- For efficient computation: for each attribute,  $Gini = 0/10 \times 1 + 10/10 \times 0.42 = 0 + 0.42 = 0.42$ 
  - Sort the attribute on values
  - Linearly scan these values, each time updating the count matrix and computing impurity
  - Choose the split position that has the least impurity

$$P(\text{yes+no}) \times Gini(\leq 55) + P(\text{yes+no}) \times Gini(> 55)$$

total no. of records

	<b>Cheat</b>	<b>No</b>		<b>No</b>		<b>No</b>		<b>Yes</b>		<b>Yes</b>		<b>Yes</b>		<b>No</b>		<b>No</b>		<b>No</b>		<b>No</b>			
		<b>Taxable Income</b>																					
Sorted Values	→	<b>60</b>		<b>70</b>		<b>75</b>		<b>85</b>		<b>90</b>		<b>95</b>		<b>100</b>		<b>120</b>		<b>125</b>		<b>220</b>			
Split Positions	→	<b>55</b>		<b>65</b>		<b>72</b>		<b>80</b>		<b>87</b>		<b>92</b>		<b>97</b>		<b>110</b>		<b>122</b>		<b>172</b>		<b>230</b>	
		<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>	<b>&lt;=</b>	<b>&gt;</b>
	<b>Yes</b>	0	3	0	3	0	3	0	3	1	2	2	1	3	0	3	0	3	0	3	0	3	0
	<b>No</b>	0	7	1	6	2	5	3	4	3	4	3	4	3	4	4	3	5	2	6	1	7	0
	<b>Gini</b>	0.420		0.400		0.375		0.343		0.417		0.400		0.300		0.343		0.375		0.400		0.420	

# Splitting based on impurity

- Impurity measures favor attributes with large number of values
- A test condition with large number of outcomes may not be desirable
  - # of records in each partition is too small to make predictions

# Splitting based on INFO

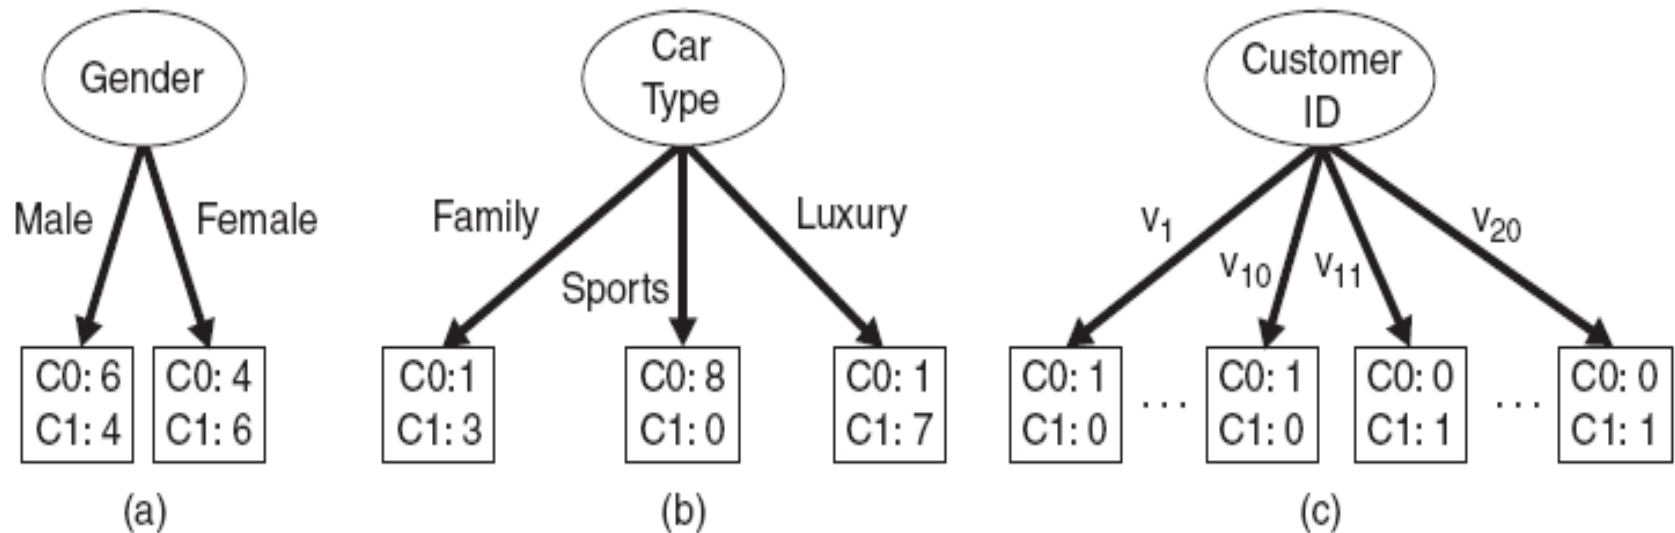


Figure 4.12. Multiway versus binary splits.

## Gain Ratio

- Splitting using information gain

$$GainRATIO_{split} = \frac{GAIN_{Split}}{SplitINFO} \quad SplitINFO = -\sum_{i=1}^k \frac{n_i}{n} \log \frac{n_i}{n}$$

Parent Node, p is split into k partitions

$n_i$  is the number of records in partition i

- Adjusts Information Gain by the entropy of the partitioning (SplitINFO). Higher entropy partitioning (large number of small partitions) is penalized!
- Used in C4.5
- Designed to overcome the disadvantage of impurity

# Stopping Criteria for Tree Induction

- Stop expanding a node when all the records belong to the same class
- Stop expanding a node when all the records have similar attribute values
- **Early termination** (to be discussed later)

# Decision Tree Based Classification

- Advantages:
  - Inexpensive to construct
  - Extremely fast at classifying unknown records
  - Easy to interpret for small-sized trees
  - Accuracy is comparable to other classification techniques for many simple data sets

# Example: C4.5


- Simple depth-first construction.
- Uses Information Gain
- Sorts Continuous Attributes at each node.
- Needs entire data to fit in memory.
- Unsuitable for Large Datasets.
  - Needs out-of-core sorting.
- You can download the software from:  
<http://www.cse.unsw.edu.au/~quinlan/c4.5r8.tar.gz>

# Other Issues

- Data Fragmentation
- Expressiveness



# Data Fragmentation




- Number of instances gets smaller as you traverse down the tree
- Number of instances at the leaf nodes could be too small to make any statistically significant decision
- *solution* You can introduce a lower bound on the number of items per leaf node in the stopping criterion.


# Expressiveness

- A classifier defines a **function** that discriminates between two (or more) classes.
- The **expressiveness** of a classifier is the **class of functions** that it can model, and the kind of data that it can **separate**
  - When we have **discrete** (or binary) values, we are interested in the class of **boolean functions** that can be modeled
  - If the data-points are real vectors we talk about the **decision boundary** that the classifier can model

# Decision Boundary



*limit your expressiveness of your model*





- Border line between two neighboring regions of different classes is known as **decision boundary**
- Decision boundary is **parallel to axes** because test condition involves a single attribute at-a-time

# Expressiveness

- Decision tree provides **expressive** representation for learning discrete-valued function
  - But they do not generalize well to certain types of Boolean functions
    - Example: **parity function**:
      - Class = 1 if there is an **even** number of Boolean attributes with truth value = True
      - Class = 0 if there is an **odd** number of Boolean attributes with truth value = True
    - For accurate modeling, must have a complete tree
- Less expressive for modeling continuous variables
  - Particularly when test condition involves only a single attribute at-a-time

# Oblique Decision Trees




- Test condition may involve multiple attributes
- More expressive representation
- Finding optimal test condition is computationally expensive

# Practical Issues of Classification

- Underfitting and Overfitting
- Evaluation

# Underfitting and Overfitting (Example)



500 circular and 500 triangular data points.


Circular points:

$$0.5 \leq \sqrt{x_1^2 + x_2^2} \leq 1$$

Triangular points:

$$\sqrt{x_1^2 + x_2^2} > 0.5 \text{ or } \sqrt{x_1^2 + x_2^2} < 1$$

# Underfitting and Overfitting




**Underfitting:** when model is **too simple**, both training and test errors are large

**Overfitting:** when model is **too complex** it models the details of the training set and fails on the test set




# Overfitting due to Noise



Decision boundary is distorted by noise point

# Overfitting due to Insufficient Examples



Lack of data points in the lower half of the diagram makes it difficult to predict correctly the class labels of that region

- Insufficient number of training records in the region causes the decision tree to predict the test examples using other training records that are irrelevant to the classification task

# Notes on Overfitting

- Overfitting results in decision trees that are more complex than necessary
- Training error no longer provides a good estimate of how well the tree will perform on previously unseen records
  - The model does not **generalize** well
- Need new ways for estimating errors

$$\Omega(t) = 0.5; \quad \Omega(T) = 1$$

# Estimating Generalization Errors

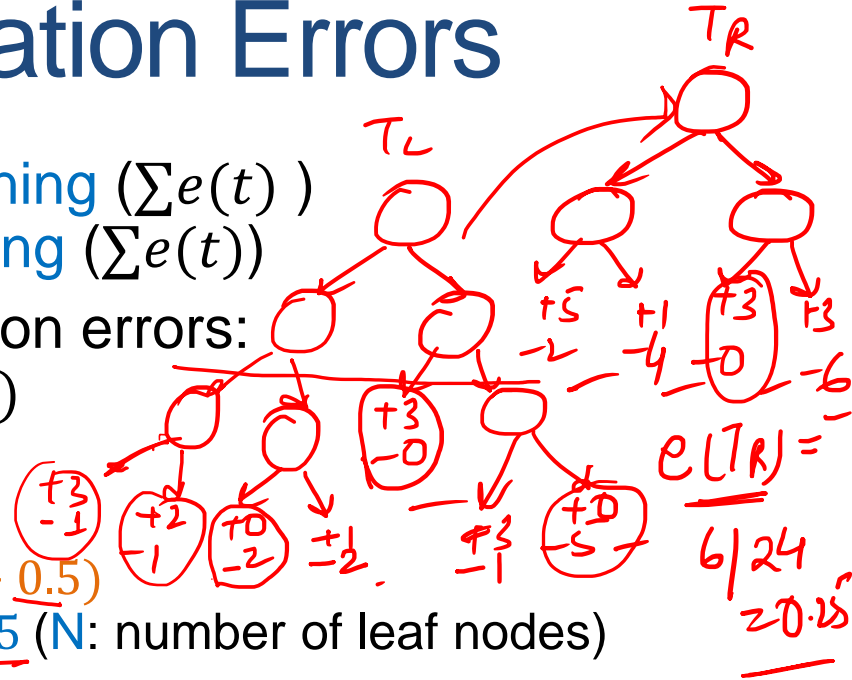
- **Re-substitution errors:** error on **training** ( $\sum e(t)$ )
- **Generalization errors:** error on **testing** ( $\sum e(t)$ )
- Methods for estimating generalization errors:
  - **Optimistic approach:**  $e'(t) = e(t)$

- **Pessimistic approach:**

- For each leaf node:  $e'(t) = (e(t) + 0.5)$
- Total errors:  $e'(T) = e(T) + N \times 0.5$  (N: number of leaf nodes)
  - Penalize large trees
- For a tree with 30 leaf nodes and 10 errors on training (out of 1000 instances)
  - Training error =  $10/1000 = 1\%$
  - Generalization error =  $(10 + 30 \times 0.5)/1000 = 2.5\%$

## Using validation set:

- Split data into **training**, **validation**, **test**
- Use **validation dataset** to estimate generalization error
- Drawback: less data for training.



$$e(T_L) = 4/24 = 0.167$$

$$e(T_R) = 4/24 = 0.167$$

$$e(T) = 10/1000 = 0.01$$

$$e'(T) = (10 + 30 \times 0.5)/1000 = 0.025$$

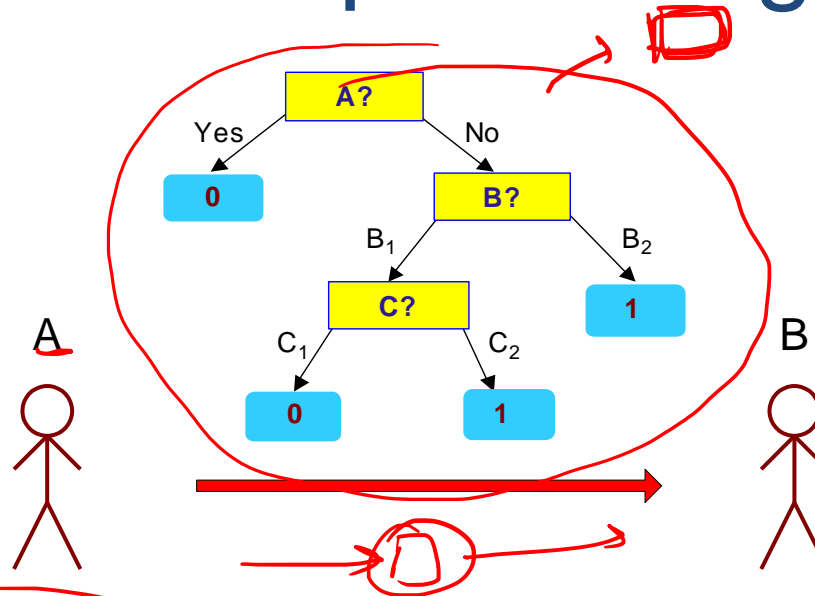
# Occam's Razor *(principle of parsimony)*

- Given two models of similar generalization errors, one should prefer the simpler model over the more complex model
- For complex models, there is a greater chance that it was fitted accidentally by errors in data
- Therefore, one should include model complexity when evaluating a model

# Minimum Description Length (MDL)

 $\checkmark \quad \checkmark$   

X	y
$X_1$	1
$X_2$	0
$X_3$	0
$X_4$	1
...	...
$X_n$	1



$\Theta(n)$


X	y
$X_1$	?
$X_2$	?
$X_3$	?
$X_4$	?
...	...
$X_n$	?

- $\text{Cost}(\text{Model}, \text{Data}) = \text{Cost}(\text{Data} | \text{Model}) + \text{Cost}(\text{Model})$ 
    - Search for the least costly model.
  - $\text{Cost}(\text{Data} | \text{Model})$  encodes the misclassification errors.
  - $\text{Cost}(\text{Model})$  encodes the decision tree
    - node encoding (number of children) plus splitting condition encoding.
- less complex models*

# How to Address Overfitting

- Pre-Pruning (Early Stopping Rule)

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
  - Stop if all instances belong to the same class
  - Stop if all the attribute values are the same



- More restrictive conditions:

- Stop if number of instances is less than some user-specified threshold
- Stop if class distribution of instances are independent of the available features (e.g., using  $\chi^2$  test)
- Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

# How to Address Overfitting...

- **Post-pruning**
  - Grow decision tree to its entirety
  - Trim the nodes of the decision tree in a **bottom-up** fashion
  - If generalization error improves after trimming, replace sub-tree by a leaf node.
  - Class label of leaf node is determined from majority class of instances in the sub-tree
- Can use **MDL** for post-pruning



# Example of Post-Pruning

Class = Yes	20
Class = No	<u>10</u>
Error = <u>10/30</u>	

Training Error (Before splitting) = 10/30

Pessimistic error =  $(10 + 0.5)/30 = 10.5/30$

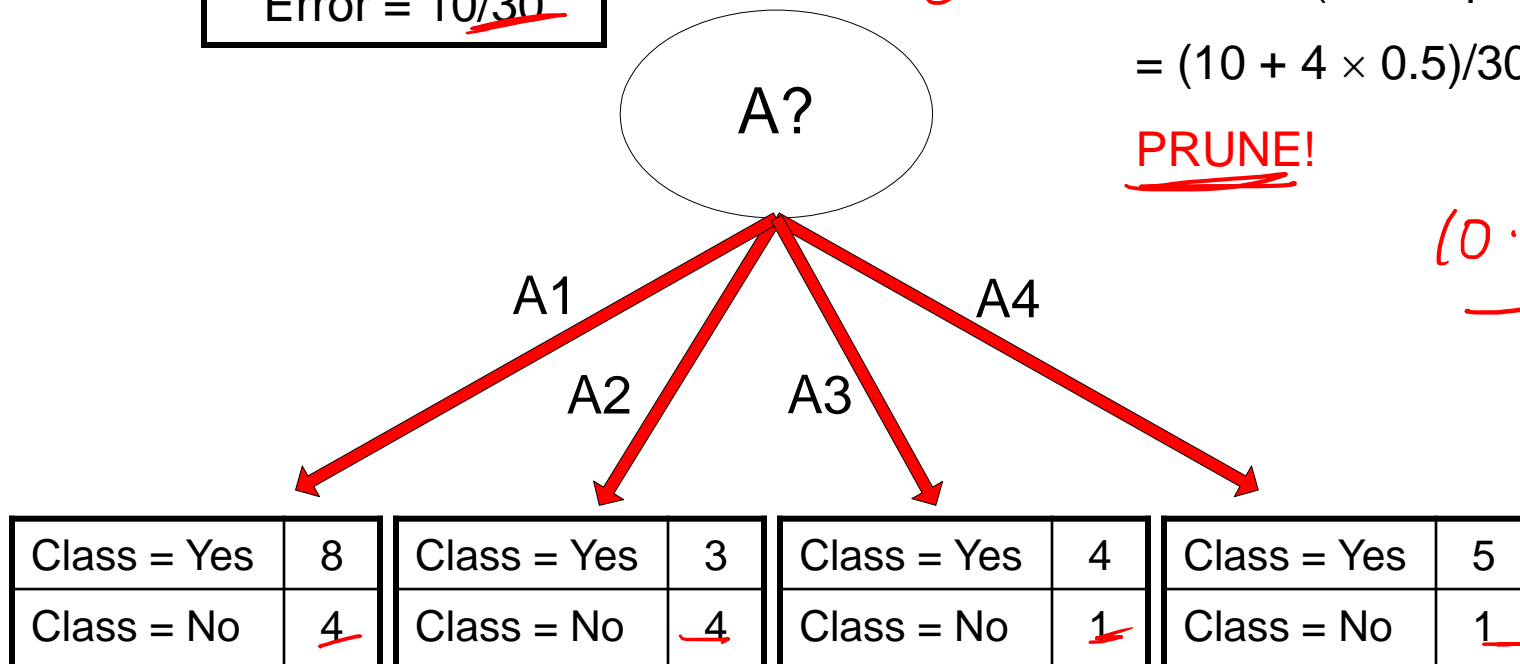
Training Error (After splitting) = 9/30 10/30

Pessimistic error (After splitting)

=  $(10 + 4 \times 0.5)/30 = \underline{11/30}$

PRUNE!

$10.5 \rightarrow 11$



# Model Evaluation

- Metrics for Performance Evaluation
  - How to evaluate the performance of a model?
- Methods for Performance Evaluation
  - How to obtain reliable estimates?
- Methods for Model Comparison
  - How to compare the relative performance among competing models?

# Model Evaluation

- Metrics for Performance Evaluation
  - How to evaluate the performance of a model?
- Methods for Performance Evaluation
  - How to obtain reliable estimates?
- Methods for Model Comparison
  - How to compare the relative performance among competing models?

# Metrics for Performance Evaluation

- Focus on the **predictive capability** of a model
  - Rather than how fast it takes to classify or build models, scalability, etc.
- **Confusion Matrix:**

ACTUAL CLASS	PREDICTED CLASS	
	Class=Yes	Class=No
Class=Yes	<b>a</b>	<b>b</b>
	<b>c</b>	<b>d</b>

a: TP (true positive)

b: FN (false negative)

c: FP (false positive)

d: TN (true negative)

# Metrics for Performance Evaluation...

	PREDICTED CLASS		
		Class=Yes	Class=No
	ACTUAL CLASS		
	Class=Yes	a (TP)	b (FN)
	Class=No	c (FP)	d (TN)

- Most widely-used metric:

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

# Limitation of Accuracy

- Consider a 2-class problem
    - Number of Class 0 examples = 9990
    - Number of Class 1 examples = 10
  - If model predicts everything to be class 0, accuracy is  $9990/10000 = \underline{99.9\%}$ 
    - Accuracy is misleading because model does not detect any class 1 example
- Handwritten notes in red:*  
w<sub>1</sub> ← lower weight  
w<sub>0</sub> ← higher weight  
w<sub>2</sub>

# Cost Matrix

	PREDICTED CLASS		
	$C(i j)$	Class=Yes	Class=No
	ACTUAL CLASS	<i>a</i>	<i>b</i>
		$C(\text{Yes} \text{Yes})$ $p(\text{yes} \text{yes})$	$C(\text{No} \text{Yes})$ $p(\text{No} \text{yes})$
		$C(\text{Yes} \text{No})$ $p(\text{yes} \text{No})$	$C(\text{No} \text{No})$ $p(\text{No} \text{No})$

**$C(i|j)$** : Cost of classifying class **j** example as class **i**

$$\text{Weighted Accuracy} = \frac{w_1 a + w_4 d}{w_1 a + w_2 b + w_3 c + w_4 d}$$

# Computing Cost of Classification

Cost Matrix	PREDICTED CLASS		
ACTUAL CLASS	C(i j)	+	-
	+	-1	100
	-	1	0

Model $M_1$	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	150	40
	-	60	250

Accuracy = 80%

Cost = 3910



Model $M_2$	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	250	45
	-	5	200

Accuracy = 90%

Cost = 4255





# Cost vs Accuracy

Count	PREDICTED CLASS		
ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	a	b
	Class=No	c	d

Accuracy is proportional to cost if

1.  $C(\text{Yes}|\text{No})=C(\text{No}|\text{Yes}) = q$
2.  $C(\text{Yes}|\text{Yes})=C(\text{No}|\text{No}) = p$

$$N = a + b + c + d$$

$$\text{Accuracy} = (a + d)/N$$

Cost	PREDICTED CLASS		
ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	p	q
	Class=No	q	p

$$\text{Cost} = p (a + d) + q (b + c)$$

$$= p (a + d) + q (N - a - d)$$

$$= q N - (q - p)(a + d)$$

$$= N [q - (q-p) \times \text{Accuracy}]$$

# Precision-Recall

rate of true results  
out of the total +ve results detected  
by the model

$$\text{Precision (p)} = \frac{a}{a+c} = \frac{TP}{TP+FP}$$

$$\text{Recall (r)} = \frac{a}{a+b} = \frac{TP}{TP+FN}$$

$$\text{F-measure (F)} = \frac{1}{\left(\frac{1/r + 1/p}{2}\right)} = \frac{2rp}{r+p} = \frac{2a}{2a+b+c} = \frac{2TP}{2TP+FP+FN}$$

total +ve  
actually  
results


Count	PREDICTED CLASS		
		Class=Yes	Class=No
	Class=Yes	a	b
ACTUAL CLASS	Class=No	c	d

$$\frac{2 \times P \times R}{P+R}$$

- Precision is biased towards C(Yes|Yes) & C(Yes|No)
- Recall is biased towards C(Yes|Yes) & C(No|Yes)
- F-measure is biased towards all except C(No|No)

# Precision-Recall plot

- Usually for parameterized models, it controls the precision/recall tradeoff



# Model Evaluation

- Metrics for Performance Evaluation
  - How to evaluate the performance of a model?
- Methods for Performance Evaluation
  - How to obtain reliable estimates?
- Methods for Model Comparison
  - How to compare the relative performance among competing models?

# Methods for Performance Evaluation

- How to obtain a reliable estimate of performance?
- Performance of a model may depend on other factors besides the learning algorithm:
  - Class distribution ✓
  - Cost of misclassification ✓
  - Size of training and test sets ✓

# Methods of Estimation

- Holdout

- Reserve  $\frac{2}{3}$  for training and  $\frac{1}{3}$  for testing

- Random subsampling


- One sample may be biased -- Repeated holdout

- Cross validation


- Partition data into  $k$  disjoint subsets
- $k$ -fold: train on  $k-1$  partitions, test on the remaining one
- Leave-one-out:  $k=n$
- Guarantees that each record is used the same number of times for training and testing

- Bootstrap

- Sampling with replacement
- ~63% of records used for training, ~27% for testing




# Dealing with class Imbalance

- If the class we are interested in is very rare, then the classifier will ignore it.
    - The class imbalance problem
  - Solution
    - We can modify the optimization criterion by using a cost sensitive metric
    - We can **balance** the class distribution
      - Sample from the larger class so that the size of the two classes is the same
      - Replicate the data of the class of interest so that the classes are balanced
        - Over-fitting issues
- 

Handwritten notes showing a sequence of operations:

$10 + 2 \rightarrow 10$  (with a red arrow pointing to the 10)  
 $10 - 2 \rightarrow 8$  (with a red arrow pointing to the 10)  
 $10 - 2 \rightarrow 8$  (with a red arrow pointing to the 10)

# Learning Curve



□ Learning curve shows how accuracy changes with varying sample size

□ Requires a sampling schedule for creating learning curve

Effect of small sample size:

- Bias in the estimate
- Variance of estimate



# Model Evaluation

- Metrics for Performance Evaluation
  - How to evaluate the performance of a model?
- Methods for Performance Evaluation
  - How to obtain reliable estimates?
- **Methods for Model Comparison**
  - How to compare the relative performance among competing models?

# ROC (Receiver Operating Characteristic)

- Developed in 1950s for signal detection theory to analyze noisy signals
  - Characterize the trade-off between positive hits and false alarms
- **ROC** curve plots **TPR** (on the **y**-axis) against **FPR** (on the **x**-axis)

$$\text{TPR} = \frac{TP}{TP + FN} = \text{Recall}$$

Fraction of **positive instances** predicted **correctly**

$$\text{FPR} = \frac{FP}{FP + TN}$$

Fraction of **negative instances** predicted **incorrectly**

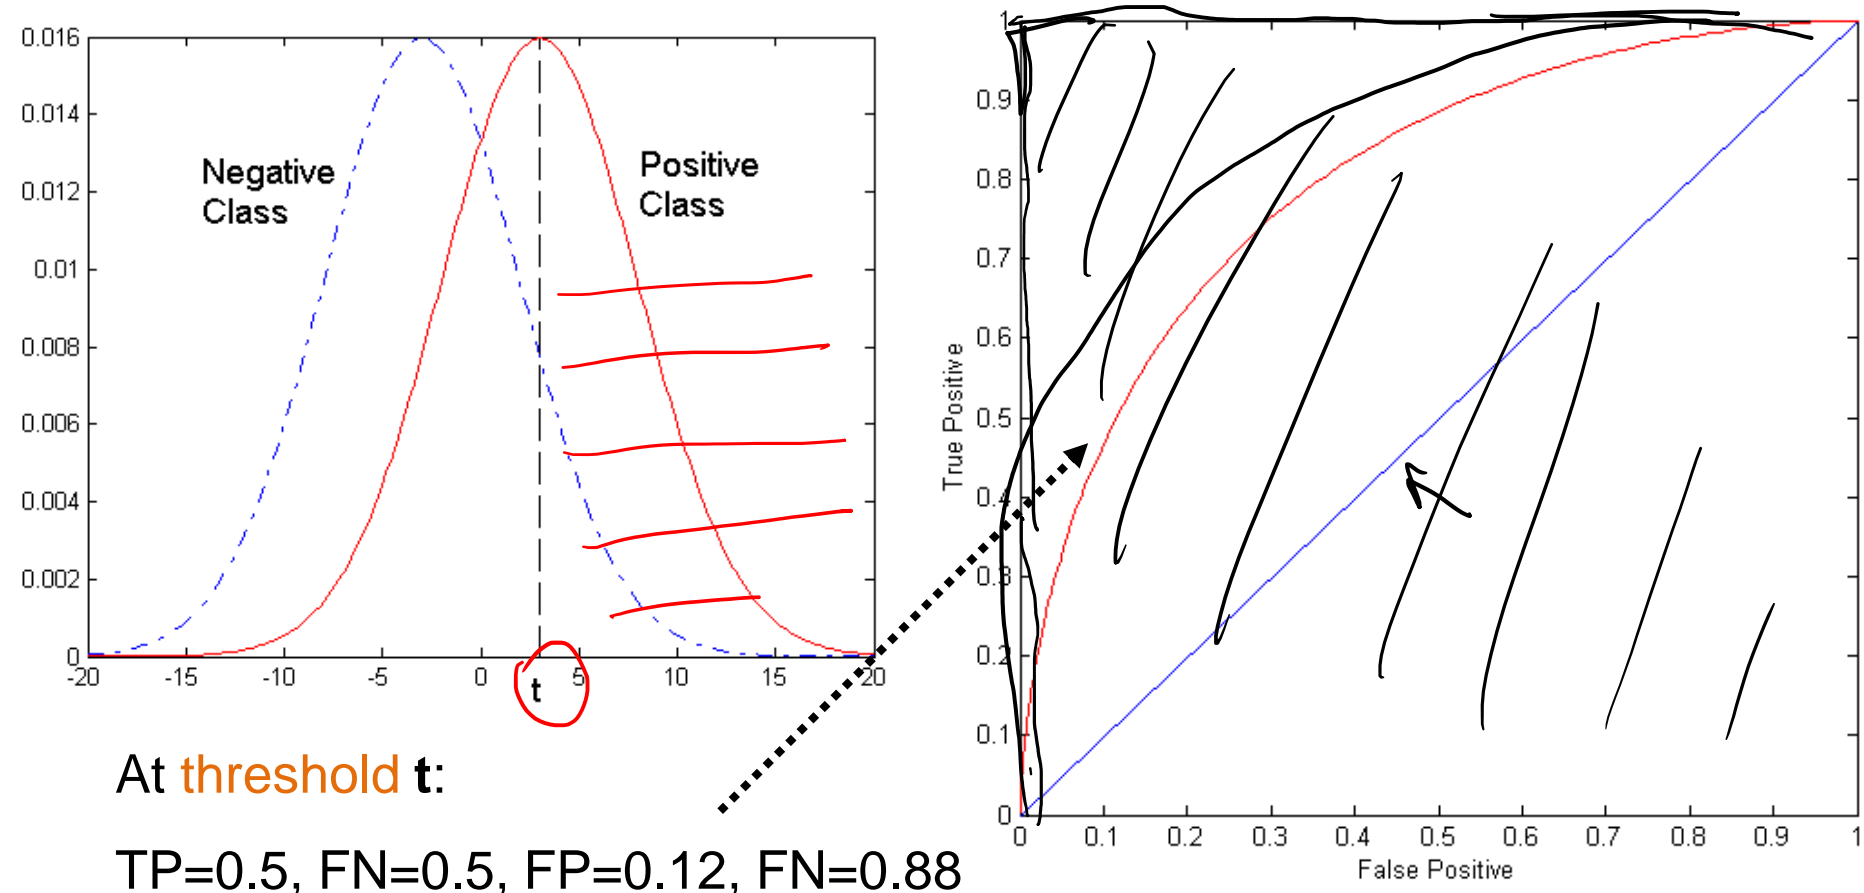
	PREDICTED CLASS		
		Yes	No
Actual	Yes	a (TP)	b (FN)
	No	c (FP)	d (TN)

# ROC (Receiver Operating Characteristic)

- Performance of a classifier represented as a **point** on the **ROC** curve
- Changing some **parameter** of the algorithm, **sample** distribution or **cost matrix** changes the location of the point

# ROC Curve


- **1**-dimensional data set containing **2** classes (**positive** and **negative**)
- any points located at  $x > t$  is classified as **positive**



# ROC Curve


(TP,FP):

- (0,0): declare everything to be negative class
  - (1,1): declare everything to be positive class
  - (1,0): ideal
- 
- Diagonal line:
    - Random guessing
    - Below diagonal line:
      - prediction is opposite of the true class





		PREDICTED CLASS	
		Yes	No
Actual	Yes	a (TP)	b (FN)
	No	c (FP)	d (TN)

# Using ROC for Model Comparison



- No model consistently outperform the other
  - $M_1$  is better for small FPR
  - $M_2$  is better for large FPR
- Area Under the ROC curve (**AUC**)
  - Ideal: Area = 1
  - Random guess:
    - Area = 0.5

# ROC curve vs Precision-Recall curve



Area Under the Curve (AUC) as a single number for evaluation